

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from xgboost import XGBClassifier
from sklearn import metrics

import warnings
warnings.filterwarnings('ignore')

from google.colab import files
uploaded = files.upload()

import io
df = pd.read_csv(io.BytesIO(uploaded['google.csv']))
print(df.head())
print('\n', df)
from sklearn.preprocessing import MinMaxScaler
from keras.layers import LSTM, Dense, Dropout
from sklearn.model_selection import TimeSeriesSplit
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.dates as mandates
from sklearn.preprocessing import MinMaxScaler
from sklearn import linear_model
from keras.models import Sequential
from keras.layers import Dense
import keras.backend as K
from keras.callbacks import EarlyStopping
from tensorflow.keras.optimizers import Adam
from keras.models import load_model
from keras.layers import LSTM
from keras.utils.vis_utils import plot_model
```



Choose Files google.csv

- **google.csv**(text/csv) - 158056 bytes, last modified: 11/10/2022 - 100% done

Saving google.csv to google.csv

|   | Date      | Open   | High   | Low    | Close  | Volume  | Year |
|---|-----------|--------|--------|--------|--------|---------|------|
| 0 | 30-Jun-17 | 943.99 | 945.00 | 929.61 | 929.68 | 2287662 | 2017 |
| 1 | 29-Jun-17 | 951.35 | 951.66 | 929.60 | 937.82 | 3206674 | 2017 |
| 2 | 28-Jun-17 | 950.66 | 963.24 | 936.16 | 961.01 | 2745568 | 2017 |
| 3 | 27-Jun-17 | 961.60 | 967.22 | 947.09 | 948.09 | 2443602 | 2017 |

df.shape

(3145, 7)

1 29-Jun-17 951.35 951.66 929.60 937.82 3206674 2017

df.describe()

|              | Open        | High        | Low         | Close       | Volume       | Year        |
|--------------|-------------|-------------|-------------|-------------|--------------|-------------|
| <b>count</b> | 3145.000000 | 3145.000000 | 3145.000000 | 3145.000000 | 3.145000e+03 | 3145.000000 |
| <b>mean</b>  | 382.514169  | 385.872099  | 378.737126  | 382.350248  | 4.205708e+06 | 2010.759300 |
| <b>std</b>   | 213.520466  | 214.636421  | 212.113835  | 213.469899  | 3.878100e+06 | 3.614485    |
| <b>min</b>   | 87.740000   | 89.290000   | 86.370000   | 87.580000   | 5.211410e+05 | 2005.000000 |
| <b>25%</b>   | 232.380000  | 234.890000  | 230.400000  | 232.440000  | 1.889613e+06 | 2008.000000 |
| <b>50%</b>   | 296.280000  | 298.520000  | 293.640000  | 296.050000  | 2.811069e+06 | 2011.000000 |
| <b>75%</b>   | 544.000000  | 548.220000  | 539.850000  | 543.650000  | 5.232088e+06 | 2014.000000 |
| <b>max</b>   | 1005.490000 | 1008.610000 | 996.620000  | 1004.280000 | 4.118289e+07 | 2017.000000 |

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3145 entries, 0 to 3144
Data columns (total 7 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Date    3145 non-null    object
1    Open     3145 non-null    float64
2    High     3145 non-null    float64
3    Low      3145 non-null    float64
4    Close    3145 non-null    float64
5    Volume   3145 non-null    int64
6    Year      3145 non-null    int64
dtypes: float64(4), int64(2), object(1)
memory usage: 172.1+ KB
```

```
plt.figure(figsize=(20,10))
plt.plot(df['Close'])
plt.title('Google Close price.', fontsize=18)
plt.ylabel('Price in dollars.')
plt.show()
```



```
df.head()
```

|   | Date      | Open   | High   | Low    | Close  | Volume  | Year |
|---|-----------|--------|--------|--------|--------|---------|------|
| 0 | 30-Jun-17 | 943.99 | 945.00 | 929.61 | 929.68 | 2287662 | 2017 |
| 1 | 29-Jun-17 | 951.35 | 951.66 | 929.60 | 937.82 | 3206674 | 2017 |
| 2 | 28-Jun-17 | 950.66 | 963.24 | 936.16 | 961.01 | 2745568 | 2017 |
| 3 | 27-Jun-17 | 961.60 | 967.22 | 947.09 | 948.09 | 2443602 | 2017 |
| 4 | 26-Jun-17 | 990.00 | 993.99 | 970.33 | 972.09 | 1517912 | 2017 |

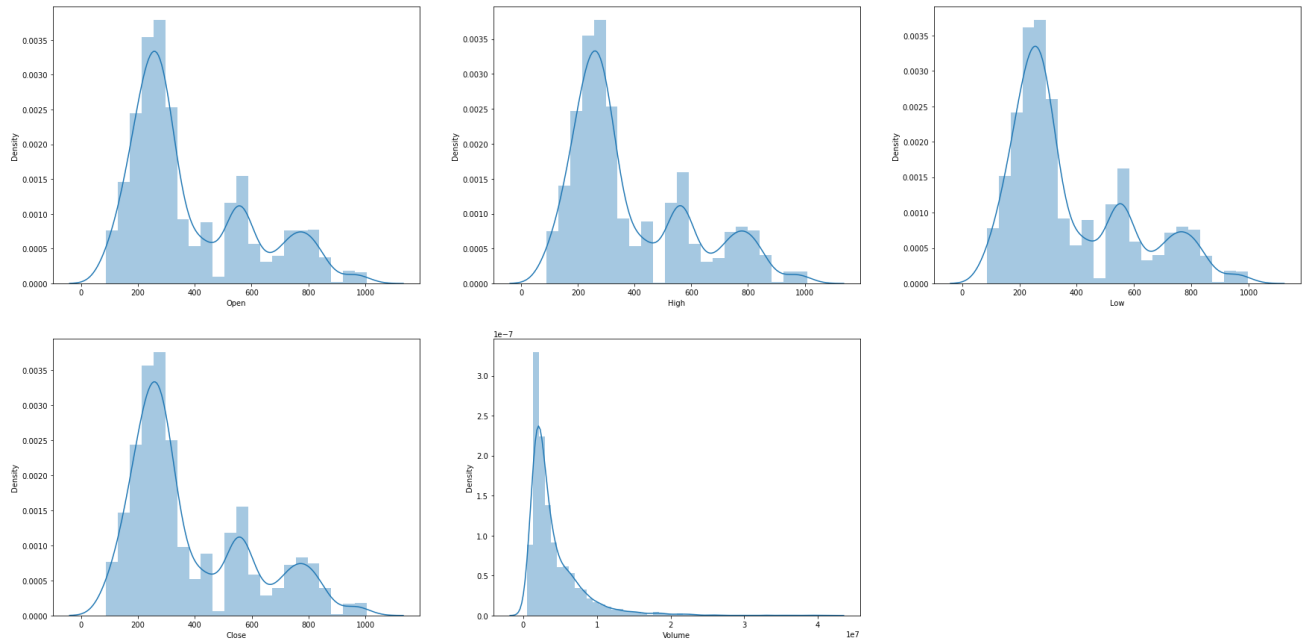
```
df.isnull().sum()
```

```
Date      0
Open      0
High      0
Low       0
Close     0
Volume    0
Year      0
dtype: int64
```

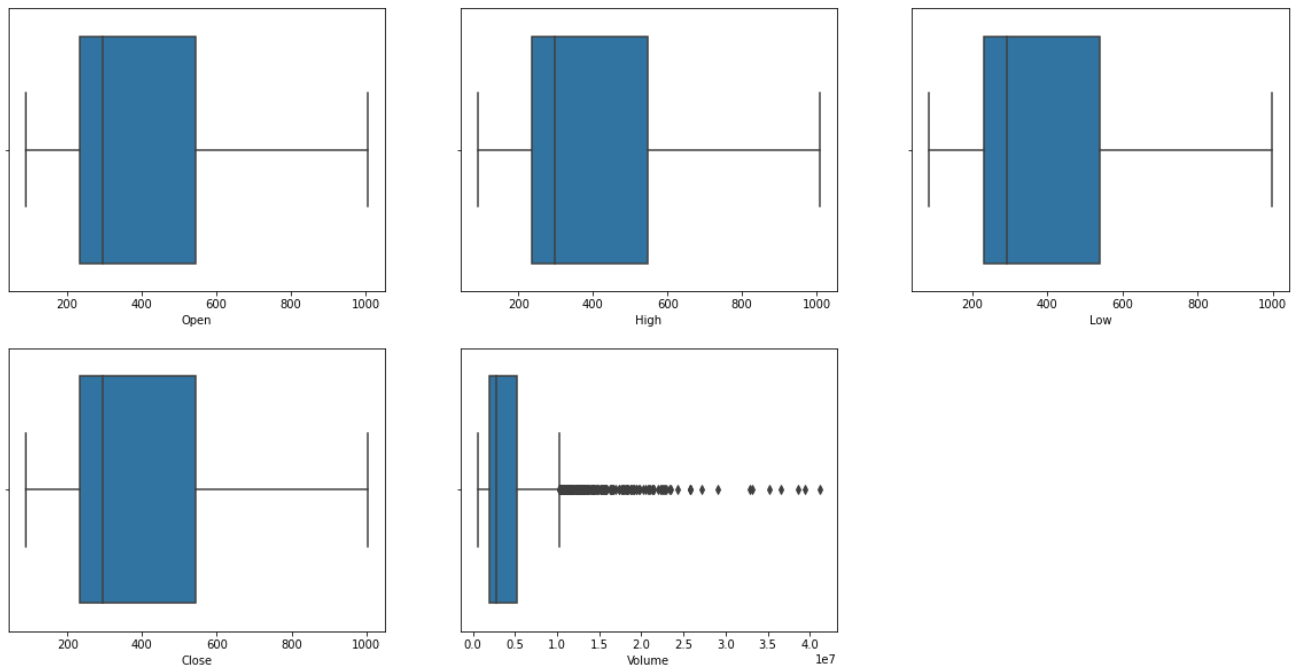
```
features = ['Open', 'High', 'Low', 'Close', 'Volume']
```

```
plt.subplots(figsize=(30,15))
```


```
for i, col in enumerate(features):
    plt.subplot(2,3,i+1)
    sb.distplot(df[col])
plt.show()
```



```
plt.subplots(figsize=(20,10))
for i, col in enumerate(features):
    plt.subplot(2,3,i+1)
    sb.boxplot(df[col])
plt.show()
```




```
df['open-close'] = df['Open'] - df['Close']
df['low-high'] = df['Low'] - df['High']
df['target'] = np.where(df['Close'].shift(-1) > df['Close'], 1, 0)
plt.pie(df['target'].value_counts().values,
        labels=[0, 1], autopct='%1.1f%%')
plt.show()
```




```
plt.figure(figsize=(10, 10))
```

```
sb.heatmap(df.corr() > 0.9, annot=True, cbar=False)
plt.show()
```



```
df['Close'].plot()
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fabef61b450>



#Scaling

```

from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
feature_transform = scaler.fit_transform(df[features])
feature_transform= pd.DataFrame(columns=features, data=feature_transform, index=df.index)
feature_transform.head()

```

|   | Open     | High     | Low      | Close    | Volume   |
|---|----------|----------|----------|----------|----------|
| 0 | 0.932988 | 0.930808 | 0.926383 | 0.918621 | 0.043444 |
| 1 | 0.941008 | 0.938052 | 0.926372 | 0.927501 | 0.066046 |
| 2 | 0.940256 | 0.950648 | 0.933579 | 0.952798 | 0.054706 |
| 3 | 0.952177 | 0.954978 | 0.945586 | 0.938704 | 0.047279 |
| 4 | 0.983122 | 0.984097 | 0.971118 | 0.964885 | 0.024514 |

```

#Set Target Variable
output_var = pd.DataFrame(df['Close'])
#Selecting the Features
features = ['Open', 'High', 'Low', 'Volume']

#Splitting to Training set and Test set
timesplit= TimeSeriesSplit(n_splits=10)
for train_index, test_index in timesplit.split(feature_transform):
    X_train, X_test = feature_transform[:len(train_index)], feature_transform[len(trai
    y_train, y_test = output_var[:len(train_index)].values.ravel(), output_var[len(trai

#Process the data for LSTM
trainX =np.array(X_train)
testX =np.array(X_test)
X_train = trainX.reshape(X_train.shape[0], 1, X_train.shape[1])
X_test = testX.reshape(X_test.shape[0], 1, X_test.shape[1])

#Building the LSTM Model
lstm = Sequential()
lstm.add(LSTM(32, input_shape=(1, trainX.shape[1]), activation='relu', return_sequences=Fa
lstm.add(Dense(1))
lstm.compile(loss='mean_squared_error', optimizer='adam')
plot_model(lstm, show_shapes=True, show_layer_names=True)

```

|            |         |                |
|------------|---------|----------------|
| lstm_input | input:  | [(None, 1, 5)] |
| InputLayer | output: | [(None, 1, 5)] |

```
history=lstm.fit(X_train, y_train, epochs=100, batch_size=8, verbose=1, shuffle=False)
```

```
Epoch 72/100
358/358 [=====] - 1s 3ms/step - loss: 15.1201
Epoch 73/100
358/358 [=====] - 1s 3ms/step - loss: 15.0143
Epoch 74/100
358/358 [=====] - 1s 3ms/step - loss: 14.9083
Epoch 75/100
358/358 [=====] - 1s 3ms/step - loss: 14.8014
Epoch 76/100
358/358 [=====] - 1s 3ms/step - loss: 14.6936
Epoch 77/100
358/358 [=====] - 1s 3ms/step - loss: 14.5846
Epoch 78/100
358/358 [=====] - 1s 3ms/step - loss: 14.4744
Epoch 79/100
358/358 [=====] - 1s 3ms/step - loss: 14.3626
Epoch 80/100
358/358 [=====] - 1s 3ms/step - loss: 14.2493
Epoch 81/100
358/358 [=====] - 1s 3ms/step - loss: 14.1345
Epoch 82/100
358/358 [=====] - 1s 3ms/step - loss: 14.0179
Epoch 83/100
358/358 [=====] - 1s 3ms/step - loss: 13.8997
Epoch 84/100
358/358 [=====] - 1s 3ms/step - loss: 13.7798
Epoch 85/100
358/358 [=====] - 1s 3ms/step - loss: 13.6583
Epoch 86/100
358/358 [=====] - 1s 3ms/step - loss: 13.5352
Epoch 87/100
358/358 [=====] - 1s 3ms/step - loss: 13.4107
Epoch 88/100
358/358 [=====] - 1s 3ms/step - loss: 13.2848
Epoch 89/100
358/358 [=====] - 1s 3ms/step - loss: 13.1576
Epoch 90/100
358/358 [=====] - 1s 3ms/step - loss: 13.0291
Epoch 91/100
358/358 [=====] - 1s 3ms/step - loss: 12.8996
Epoch 92/100
358/358 [=====] - 1s 3ms/step - loss: 12.7691
Epoch 93/100
358/358 [=====] - 1s 3ms/step - loss: 12.6377
Epoch 94/100
358/358 [=====] - 1s 3ms/step - loss: 12.5056
Epoch 95/100
358/358 [=====] - 1s 3ms/step - loss: 12.3730
Epoch 96/100
358/358 [=====] - 1s 3ms/step - loss: 12.2398
Epoch 97/100
358/358 [=====] - 1s 3ms/step - loss: 12.1063
```



```
Epoch 98/100
358/358 [=====] - 1s 3ms/step - loss: 11.9726
Epoch 99/100
358/358 [=====] - 1s 3ms/step - loss: 11.8388
Epoch 100/100
358/358 [=====] - 1s 3ms/step - loss: 11.7050
```

```
#LSTM Prediction
```

```
y_pred= lstm.predict(X_test)
```

```
9/9 [=====] - 0s 2ms/step
```

```
#Predicted vs True Adj Close Value - LSTM
```

```
plt.plot(y_test, label='True Value')
```

```
plt.plot(y_pred, label='LSTM Value')
```


```
plt.title("Prediction by LSTM")
```

```
plt.xlabel('Time Scale')
```

```
plt.ylabel('Scaled USD')
```

```
plt.legend()
```

```
plt.show()
```



```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.svm import SVC
```

```
from xgboost import XGBClassifier
```

```
from sklearn import metrics
```

```
df['open-close'] = df['Open'] - df['Close']
```

```
df['low-high'] = df['Low'] - df['High']
```

```
df['target'] = np.where(df['Close'].shift(-1) > df['Close'], 1, 0)
```

```
features = df[['open-close', 'low-high', 'Year']]
```

```
target = df['target']
```

```
scaler = StandardScaler()
```

```

features = scaler.fit_transform(features)

X_train, X_valid, Y_train, Y_valid = train_test_split(
    features, target, test_size=0.1, random_state=2022)
print(X_train.shape, X_valid.shape, '\n\n')
models = [LogisticRegression(), SVC(
    kernel='poly', probability=True), XGBClassifier()]

for i in range(3):
    models[i].fit(X_train, Y_train)

    print(f'{models[i]} : ')
    print('Training Accuracy : ', metrics.roc_auc_score(
        Y_train, models[i].predict_proba(X_train)[:,:1]))
    print('Validation Accuracy : ', metrics.roc_auc_score(
        Y_valid, models[i].predict_proba(X_valid)[:,:1]))
    print()

    (2830, 3) (315, 3)

```

```

LogisticRegression() :
Training Accuracy : 0.9115537167270719
Validation Accuracy : 0.9145858585858586

```

```

SVC(kernel='poly', probability=True) :
Training Accuracy : 0.8955785609133604
Validation Accuracy : 0.8862020202020202

```

```


XGBClassifier() :
Training Accuracy : 0.9366449247533783
Validation Accuracy : 0.9161212121212121

```

```

metrics.plot_confusion_matrix(models[0], X_valid, Y_valid)
plt.show()

```



[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 9:14 PM



Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.